

2.5 Quantifying parallelism

Given:

- A p -processor parallel computer
- A problem (e.g., linear system) and an algorithm to solve it (e.g., Gaussian elimination)

13 ♦ Definition: Let $T(p)$ denote the time taken by the parallel algorithm on p processors.

$$S(p) := \frac{T(1)}{T(p)} \text{ is the speed-up and}$$

$$E(p) := \frac{S(p)}{p} \text{ is the efficiency}$$

of the algorithm **w.r.t. itself**.

Let T' denote the time taken by the **fastest** known (serial) algorithm to solve the problem.

$$S'(p) := \frac{T'}{T(p)} \text{ is the true speed-up and}$$

$$E'(p) := \frac{S'(p)}{p} \text{ is the true efficiency}$$

of the parallel algorithm.

Let $\#ops(p)$ denote the number of operations ($+$, $-$, $*$, often also $/$, $\sqrt{}$).

$$R(p) = \frac{\#ops(p)}{T(p)} \quad \text{is the (compute) performance}$$

of the algorithm on p processors (in **Flop/s**: floating-point operations per second).

Note: By Definition 13,

$$\begin{aligned} S(p) &= \frac{\#ops(1)}{R(1)} \cdot \frac{R(p)}{\#ops(p)} \\ &\approx \frac{R(p)}{R(1)} \quad \text{if } \#ops(p) \approx \#ops(1) . \end{aligned}$$

More detailed analysis: Let n denote the “characteristic size” of the problem (e.g., matrix dimension for linear systems). Then

$$T(p) \rightsquigarrow T(p, n), \quad S(p) \rightsquigarrow S(p, n), \quad \text{etc.}$$

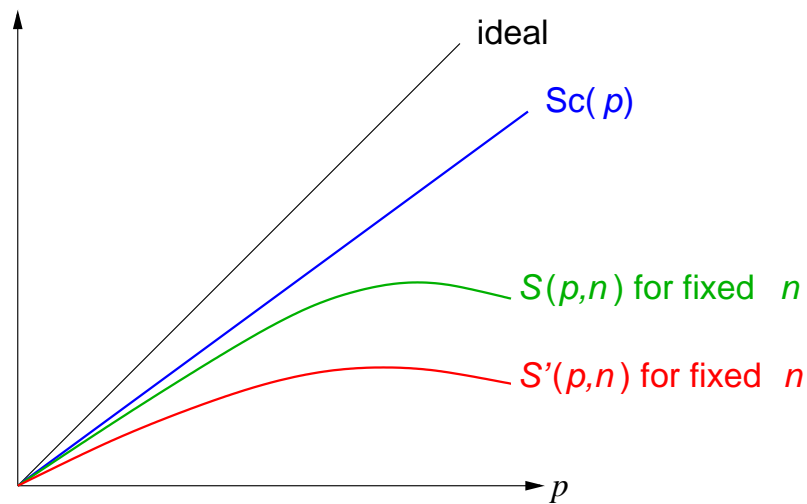
14 ♦ Definition: Given a p -processor parallel computer, the **scale-up** for a variable-size problem is given by

$$Sc(p) := \sup_{n \in \mathbb{N}} S(p, n) .$$

(In practice, n can be limited to problems that fit into the available memory.)

Note: $Sc(p) \geq S(p, n) \geq S'(p, n)$ for any n .

Typical behavior:



15 \diamond **Definition:** A parallel algorithm for solving a problem **scales** iff

$$\liminf_{p \rightarrow \infty} \frac{Sc(p)}{p} > 0$$

(i.e., the quotient $\frac{Sc(p)}{p}$ is bounded from below by some $\varepsilon > 0$).

16 ♦ Theorem: (**Amdahl's law**) Let $q \in [0, 1]$ denote the fraction of those operations in an algorithm, which cannot be parallelized. Then

$$S(p) \leq \frac{1}{q}, \quad \text{no matter how large } p \text{ is.}$$

Proof: $T(p) \geq q \cdot T(1)$. ■

17 ♦ Example: If $q = 0.01$ (i.e., 99% of all operations can be parallelized) then

$$S(p) \leq 100.$$

2.6 Communication instructions

send(D) to P_i : send the data D to processor P_i

receive(D) from P_j : receive the data D from processor P_j

broadcast(D) : send the data D to all processors (perhaps only to those of a specified group)

reduce : see Section 2.7.1

18 ♦ Remark: Sending data from one processor to another one typically involves

- setting up a connection between them (e.g., determining a suitable path through the network)
- transmitting the data

⇒ communication time for sending a length- ℓ message is

$$T_{\text{comm}}(\ell) = \alpha + \beta \cdot \ell,$$

where

- α is the **latency** or start-up time,
- β is the transmission time per byte (reciprocal of the **communication bandwidth**), and
- ℓ is the length of the message in bytes.

Realistic values: $\alpha \sim 2\mu\text{s}$, $\beta \sim 0.1\text{ns/B}$ ($\hat{=}$ 10GB/s).

Check your understanding. You should know . . .

- Speed-up $S(p)$, efficiency $E(p)$, scale-up $Sc(p)$
- What does Amdahl's law tell?
- What are the four basic communication routines?

2.7 Basic parallelization schemes

2.7.1 The fan-in principle for reduction operations

Problem: Given $a_0, \dots, a_{n-1} \in \mathbb{R}$, compute

$$s := \sum_{j=0}^{n-1} a_j .$$

19 \diamond **Algorithm:** Serial sum

$s^{(-1)} := 0$

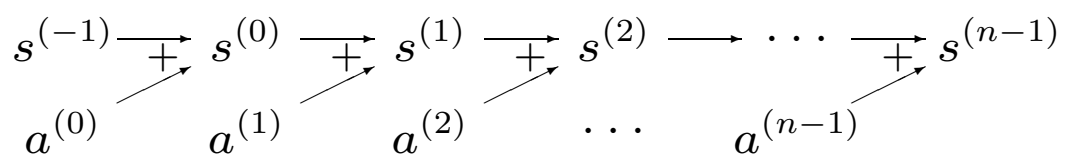
for $i = 0 : n - 1$

$\{ \text{assume that } s^{(i-1)} = \sum_{j=0}^{i-1} a_j \}$

$s^{(i)} := s^{(i-1)} + a_i$

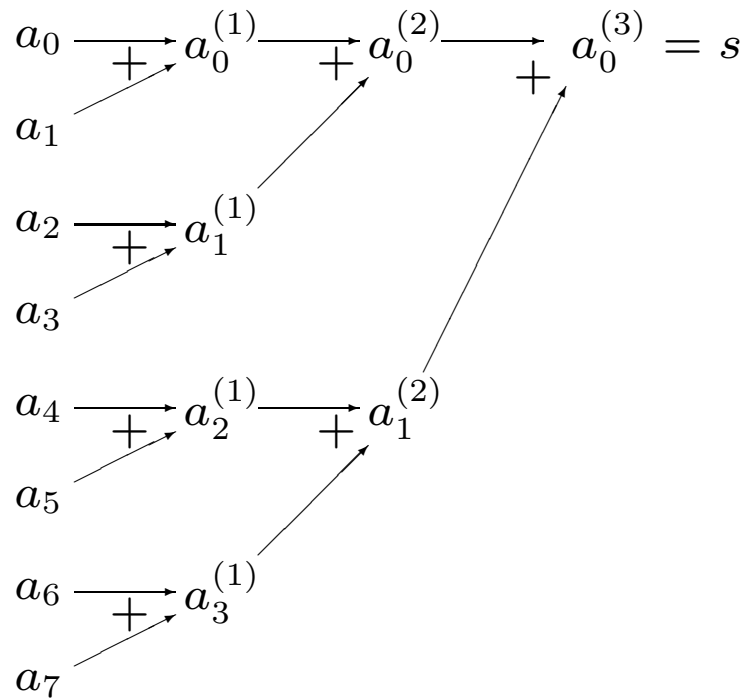
$\{ s^{(n-1)} = s \}$

Dependence graph:



Nodes represent intermediate quantities,
(directed) edges represent dependences (e.g., $s^{(0)}$ and $a^{(1)}$
must be available to compute $s^{(1)}$),
the “+” are optional, specifying the computations.

Different approach (here for $n = 8 = 2^3$):



{ to fan out: to spread out in the shape of a fan }

20 ♦ Algorithm: Serial **fan-in summation**, assumes $n = 2^N$

{ let $a_j^{(0)} \equiv a_j, j = 0 : n - 1$ }

for $k = 1 : N$

{ assume that for $i = 0 : 2^{N-(k-1)} - 1$, $a_i^{(k-1)}$ is the sum of the i th length- 2^{k-1} subsequence of a , i.e.,

$$a_i^{(k-1)} = \sum_{j=i2^{k-1}}^{(i+1)2^{k-1}-1} a_j \}$$

for $i = 0 : 2^{N-k} - 1$

$$a_i^{(k)} := a_{2i}^{(k-1)} + a_{2i+1}^{(k-1)}$$

{ $a_0^{(N)} = s$ }

Potential for parallelism: The i loop can be done in parallel.

Let $p = 2^N$, processors P_0, \dots, P_{p-1} .

For each “stage index” k , choose 2^{N-k} of the processors (labelled $P_i^{(k)}$, $i = 0 : 2^{N-k} - 1$), such that

- $P_i^{(k)}$ computes the partial sum $a_i^{(k)} := a_{2i}^{(k-1)} + a_{2i+1}^{(k-1)}$
 $\{ a_{2i}^{(k-1)}$ is available in $P_{2i}^{(k-1)}$, $a_{2i+1}^{(k-1)}$ is in $P_{2i+1}^{(k-1)}$ $\}$

21 ◇ Algorithm: Fan-in summation with $p = 2^N$ processors; all processors run the same pseudocode

$\{$ each of the 2^N processors $P_i^{(0)}$ holds $a_i =: a_{\text{local}}$ $\}$

for $k = 1 : N$

$\{$ assume that each processor $P_i^{(k-1)}$,
 $i = 0 : 2^{N-(k-1)} - 1$, holds $a_{\text{local}} \equiv a_i^{(k-1)}$ $\}$

for $i = 0 : 2^{N-k} - 1$

if $\text{me} = P_{2i}^{(k-1)}$ **or** $\text{me} = P_{2i+1}^{(k-1)}$
 $\{$ I have one summand needed for $a_i^{(k)}$ $\}$

send(a_{local}) **to** $P_i^{(k)}$

if $\text{me} = P_i^{(k)}$

$\{$ I am computing $a_i^{(k)}$ $\}$

receive(summand₁) **from** $P_{2i}^{(k-1)}$

receive(summand₂) **from** $P_{2i+1}^{(k-1)}$

$a_{\text{local}} := \text{summand}_1 + \text{summand}_2$

$\{$ $s = a_0^{(N)}$ is the result a_{local} in $P_0^{(N)}$ $\}$

22 ◇ Remark: To add $n = q \cdot 2^N$ numbers with $p = 2^N$ processors ($q, N \in \mathbb{N}$), start with the partial sums

$$a_{\text{local}} := \sum_{j=iq}^{(i+1)q-1} a_j \quad \text{in } P_i^{(0)}, \quad i = 0 : 2^N - 1.$$

Computation of a_{local} : E.g., with Algorithm 19

23 ◇ Definition: A **deadlock** in a parallel algorithm occurs when a processor wants to receive data that are never sent.

24 ◇ Remark: Depending on the implementation, even a **send** operation may terminate only when the corresponding **receive** of the partner is completed (**blocking send**).

This does not happen if outgoing data are buffered

\Rightarrow **send** completes when data are in the buffer

(unless buffer overflows).

Problem: Algorithm 21 is deadlock-free (for blocking sends) only if all $P_i^{(k-1)}$, $i = 0 : 2^{N-(k-1)} - 1$, are different from all $P_i^{(k)}$, $i = 0 : 2^{N-k} - 1$; otherwise deadlocks may occur.

This condition cannot be fulfilled for $k = 1$!

Solution: Careful selection of the processors eliminates a (critical) part of the communication.

More precisely, let

$$P_i^{(k)} := P_{i \cdot 2^k}, \quad k = 0 : N, \quad i = 0 : 2^{N-k} - 1. \quad (2.1)$$

Then

$$P_{2i}^{(k-1)} = P_{2i \cdot 2^{k-1}} = P_{i \cdot 2^k} = P_i^{(k)},$$

i.e., the **receive**(summand₁) and the first case for **send** can be skipped because the processor would send to itself.

25 ◇ Algorithm: Improved fan-in summation

```
for  $k = 1 : N$ 
  for  $i = 0 : 2^{N-k} - 1$ 
    if me =  $P_{2i+1}^{(k-1)}$  { =  $P_{i2^k+2^{k-1}}$  }
      { I have the summand  $a_{2i+1}^{(k-1)}$  that is needed
        in  $P_i^{(k)}$  to compute  $a_i^{(k)}$  }
      send(  $a_{\text{local}}$  ) to  $P_i^{(k)}$  { =  $P_{i2^k}$  }
    if me =  $P_i^{(k)}$ 
      { I am computing  $a_i^{(k)}$  }
      receive( summand2 ) from  $P_{2i+1}^{(k-1)}$ 
       $a_{\text{local}} := a_{\text{local}} + \text{summand}_2$ 
```

26 ◇ Remark: With the logical to physical mapping $P_i^{(k)} = P_{i \cdot 2^k}$, the indices of the communication partners in Algorithms 25 differ by 2^{k-1} , i.e., only in bit $k - 1$

\Rightarrow in a N -dimensional hypercube all communication is between directly connected processors.

27 ♦ Remark: In Algorithm 25, each processor performs

$$\sum_{k=1}^N \sum_{i=0}^{2^{N-k}-1} \mathcal{O}(1) = \mathcal{O}(n)$$

operations, mostly for finding out, which rôle the processor plays in stage k (in the majority of cases, none).

In practice, this must be done without a loop.

28 ♦ Algorithm: Fan-in summation

```
for  $k = 1 : N$ 
  if (me is a multiple of  $2^{k-1}$ )
    { only these procs were active in stage  $k - 1$  }
    and (me is not a multiple of  $2^k$ )
    { only these procs must send }
    send(  $a_{\text{local}}$  ) to  $P_{\text{me}-2^{k-1}}$ 
  if (me is a multiple of  $2^k$ )
    { only these procs are active in stage  $k$  }
    receive(  $\text{summand}_2$  ) from  $P_{\text{me}+2^{k-1}}$ 
     $a_{\text{local}} := a_{\text{local}} + \text{summand}_2$ 
```

\Rightarrow only $\mathcal{O}(N) = \mathcal{O}(\log n)$ operations per processor.

29 ◇ Theorem: Let $a_i \in \mathbb{R}$, $i = 0 : n - 1$ ($n = q \cdot 2^N$) and $p = 2^N$.

Neglecting communication, running Algorithm 28 to com-

pute $s := \sum_{i=0}^{n-1} a_i$ yields

$$S(p) = \frac{q \cdot 2^N - 1}{(q - 1) + N} ,$$

$$E(p) = \frac{q - \frac{1}{2^N}}{(q - 1) + N} .$$

Proof: Let T_+ denote the time for a (serial) addition

$$\Rightarrow T(1) = (n - 1) \cdot T_+ ,$$

{ $N = 0$, $q = n$ means serial summation }

$$T(p) = (q - 1) \cdot T_+$$

{ serial addition $a_i^{(0)} := \sum_{j=iq}^{(i+1)q-1} a_j$ of q
numbers in each processor $P_i^{(0)}$ }

$$+ N \cdot T_+$$

{ at most one add per processor in each
pass of the k loop }

$$\Rightarrow S(p) = \frac{T(1)}{T(p)} = \frac{q \cdot 2^N - 1}{(q - 1) + N} ,$$

$$E(p) = \frac{1}{2^N} \cdot S(p) .$$



30 ◇ Example:

$$q = 1 \Rightarrow E(p) = \frac{1 - \frac{1}{2^N}}{N} \approx \frac{1}{N} = \frac{1}{\log_2 p} \quad \{ p = 2^N \}$$

$$q \rightarrow \infty \Rightarrow E(p) \rightarrow 1 \quad \{ \text{unrealistic !} \}$$

Realistic: $q \leq C$ (**limited memory** per processor):

$$\begin{aligned} E(p) &= \frac{q - \frac{1}{2^N}}{(q - 1) + N} \\ &\leq \frac{C}{N} \\ &\rightarrow 0 \quad \text{as } p \rightarrow \infty \text{ (i.e., } N \rightarrow \infty) \end{aligned}$$

\Rightarrow the fan-in algorithm **does not scale**.

31 ◇ Remark: Since communication is expensive, the real speed-up is much lower.

With the communication model from Remark 18, the overall time for 8-byte (IEEE double precision) computations is

$$T(p) = (q - 1) \cdot T_+ + N \cdot (\alpha + 8B \cdot \beta + T_+) ,$$

where T_+ is the time for one addition, α is the start-up time for communication, and β is the transmission time per byte.

Using $q = n/2^N$ and the { JUQUEEN } values $T_+ = 1/(4 \cdot 1.6\text{GHz}) \sim 0.156\text{ns}$, $\alpha = 2.5\mu\text{s}$, and $\beta = 0.025\text{ns/B}$, one obtains

$$\begin{aligned} T(p) &= \left(\frac{n}{2^N} - 1 \right) \cdot 0.156\text{ns} \\ &\quad + N \cdot (2.5\mu\text{s} + 8\text{B} \cdot 0.025\text{ns/B} + 0.156\text{ns}) \\ &= \left(\frac{n}{2^N} - 1 \right) \cdot 0.156\text{ns} + N \cdot 2500.356\text{ns} . \end{aligned}$$

For $n = 2^{20}$,

$$T(p = 1) \approx 163.6\mu\text{s}, \quad T(p = n) \approx 50.0\mu\text{s},$$

and the time is minimized for $N = 6$, i.e., adding 1 048 576 numbers is fastest on $2^6 = 64$ processors, yielding

$$S_{\max} = S(64) = \frac{T(1)}{T(64)} \approx \frac{163.6\mu\text{s}}{17.56\mu\text{s}} \approx 9.3 .$$

Adding more processors slows down the computation.

32 ◇ Remark: The fan-in principle applies to any expression $a_0 \circ a_1 \circ \dots \circ a_{n-1}$ ($n = 2^N$), where \circ is an **associative** operation, e.g.,

- $\max_{i=0}^{n-1} a_i$ ($a \circ b = \max\{a, b\}$)
 - $\min_{i=0}^{n-1} a_i$
 - $\prod_{i=0}^{n-1} a_i$
 - $\gcd(m_0, \dots, m_{n-1})$ { greatest common divisor }
 - $\text{lcm}(m_0, \dots, m_{n-1})$ { least common multiple }
 - $\sum_{i=0}^{n-1} A_i$ { matrices }
 - $\prod_{i=0}^{n-1} A_i$
-

33 ◇ Definition: In a reduction operation,

$$\text{reduce}(a_0, \dots, a_{n-1} ; \circ ; \text{res} ; P_i, i \in I ; P_k) ,$$

the processors $P_i, i \in I$, work together to compute the result $\text{res} = a_0 \circ \dots \circ a_{n-1}$, where \circ is an associative operation.

The result is made available in P_k .

34 ♦ Theorem: (Backward stability of summation)

If the sum $s := \sum_{j=0}^{n-1} a_j$ of n floating-point numbers a_j is computed with a floating-point arithmetic satisfying

$$\underbrace{\overline{x + y}}_{\text{computed result}} = \underbrace{(x + y)}_{\text{exact result}} \cdot (1 + \delta) \quad \text{with } |\delta| \leq \varepsilon \quad (2.2)$$

for all floating-point numbers x, y and some fixed $0 < \varepsilon \ll 1$ (**machine precision**) then

$$\underbrace{\overline{s}}_{\text{computed sum}} = \sum_{j=0}^{n-1} \underbrace{\tilde{a}_j}_{\text{perturbed summands}}$$

with

$$\tilde{a}_j = a_j \prod_{\ell=j}^{n-1} (1 + \delta_\ell) \quad \text{if Algorithm 19 is used} \quad (2.3)$$

or

$$\tilde{a}_j = a_j \cdot \prod_{\nu=1}^N (1 + \delta_{j\nu}) \quad \text{for } n = 2^N \text{ and Alg. 20,} \quad (2.4)$$

where $|\delta_\ell|, |\delta_{j\nu}| \leq \varepsilon$.

35 ♦ Remarks:

- a) IEEE arithmetic guarantees (2.2), unless an under-/overflow occurs.
- b) To first order,

$$|\tilde{a}_j - a_j| \lesssim (n - j) \cdot \varepsilon \cdot |a_j| \quad \text{for Alg. 19, and}$$

$$|\tilde{a}_j - a_j| \lesssim \log_2 n \cdot \varepsilon \cdot |a_j| \quad \text{for Alg. 20 .}$$

Therefore, if no information on the magnitude of the a_j is available, Algorithm 20 gives smaller **bounds** for the errors.

Proof for Theorem 34:

For (2.3), rewrite Algorithm 19 in floating-point arithmetic:

$$\underbrace{\overline{s^{(-1)}}}_{\text{computed quantity}} := 0 \quad \{ \text{no rounding error} \}$$

for $i = 0 : n - 1$

$$\{ \text{assume } \underbrace{\overline{s^{(i-1)}}}_{\text{computed}} = \sum_{j=0}^{i-1} \underbrace{a_j}_{\text{exact inputs}} \cdot \prod_{\ell=j}^{i-1} (1 + \delta_\ell), \text{ where } |\delta_\ell| \leq \varepsilon \}$$

$$\overline{s^{(i)}} := \underbrace{\overline{s^{(i-1)}} \oplus a_i}_{\text{floating-point addition}}$$

{ Then

$$\begin{aligned}
\overline{s^{(i)}} &= (\overline{s^{(i-1)}} + a_i) \cdot (1 + \delta_i) \quad \text{for some } |\delta_i| \leq \varepsilon \quad \text{by (2.2)} \\
&= \left(\sum_{j=0}^{i-1} a_j \prod_{\ell=j}^{i-1} (1 + \delta_\ell) + a_i \right) \cdot (1 + \delta_i) \\
&= \sum_{j=0}^{i-1} a_j \prod_{\ell=j}^i (1 + \delta_\ell) + a_i (1 + \delta_i) \\
&= \sum_{j=0}^i a_j \prod_{\ell=j}^i (1 + \delta_\ell) \quad \}
\end{aligned}$$

$$\{ \text{Therefore } \overline{s} = \overline{s^{(n-1)}} = \sum_{j=0}^{n-1} a_j \prod_{\ell=j}^{n-1} (1 + \delta_\ell) \}$$

This proves (2.3).

For (2.4), rewrite Algorithm 20 in floating-point arithmetic:

$$\{ \overline{a_j^{(0)}} \equiv a_j, j = 0 : 2^N - 1 \ ; \quad \text{thus } \overline{a_j^{(0)}} = a_j \underbrace{\prod_{\nu=1}^0 (1 + \delta_{j\nu})}_{\equiv 1} \}$$

for $k = 1 : N$

$$\{ \text{assume that } \overline{a_i^{(k-1)}} = \sum_{j=i2^{k-1}}^{(i+1) \cdot 2^{k-1} - 1} a_j \prod_{\nu=1}^{k-1} (1 + \delta_{j\nu})$$

holds for $i = 0 : 2^{N-(k-1)} - 1 \}$

for $i = 0 : 2^{N-k} - 1$

$$\overline{a_i^{(k)}} := \overline{a_{2i}^{(k-1)}} \oplus \overline{a_{2i+1}^{(k-1)}}$$

{ Thus

$$\overline{a_i^{(k)}} = \left[\sum_{j=i2^k}^{i2^k+2^{k-1}-1} a_j \prod_{\nu=1}^{k-1} (1 + \delta_{j\nu}) + \sum_{j=i2^k+2^{k-1}}^{(i+1) \cdot 2^k - 1} a_j \prod_{\nu=1}^{k-1} (1 + \delta_{j\nu}) \right] \cdot (1 + \delta_{ik})$$

for some $|\delta_{ik}| \leq \varepsilon$ by (2.2)

$$= \sum_{j=i2^k}^{(i+1) \cdot 2^k - 1} a_j \prod_{\nu=1}^k (1 + \delta_{j\nu})$$

with $\delta_{jk} = \delta_{ik}$ for all $j = i2^k : (i+1)2^k - 1$

$$\{ \text{Therefore } \overline{a_0^{(N)}} = \sum_{j=0}^{2^N-1} a_j \prod_{\nu=1}^N (1 + \delta_{j\nu}) \}$$

This proves (2.4). ■

Check your understanding. You should know . . .

- How does the fan-in principle work?
- How fan-in works when $p < n$.
- Whether fan-in scales and why not.
- What a “global reduce” is.
- There is an issue regarding numerical round-off errors. What is it, exactly?
- What is a deadlock, and how can we avoid these?
- That we distinguish between blocking and non-blocking communications
- Assume that you have to compute the inner product of two vectors x and y , which are distributed over the processors. (This is a very common task!) How can we perform this using fan-in?

2.7.2 Parallel matrix multiplication

Given $A = (a_{ik}) \in \mathbb{R}^{n \times m}$ and $B = (b_{kj}) \in \mathbb{R}^{m \times q}$,

compute $C = (c_{ij}) = A \cdot B \in \mathbb{R}^{n \times q}$, where

$$c_{ij} = \sum_{k=1}^m a_{ik} b_{kj} , \quad i = 1 : n, \quad j = 1 : q .$$

Motivation: Stochastic models are often described by a **state change matrix** A , where a_{ij} is the probability that an object moves from state j to state i .

For a given initial distribution $s^{(0)}$ of the objects to the states, a new distribution $s^{(1)}$ is obtained by $s^{(1)} = A \cdot s^{(0)}$.

Analogously, for the n th distribution we have $s^{(n)} = A \cdot s^{(n-1)} = \dots = A^n \cdot s^{(0)}$, and thus the state change matrix from $s^{(0)}$ to $s^{(n)}$ is $A^n = A \cdot \dots \cdot A$.

Such models are also used in telecommunication and for search engines.

36 ♦ Definition: A **partitioning** of a set I is a system of sets, $\{ I_r : r = 1 : R \}$, such that

(i) $I_r \neq \emptyset$ for all r ,

(ii) $I_r \cap I_s = \emptyset$ for all $r \neq s$, and

(iii) $\bigcup_{r=1}^R I_r = I$.

Let the parallel computer contain $p = R \cdot S \cdot T$ processors $P_{r,s,t}$, $r = 1 : R$, $s = 1 : S$, $t = 1 : T$, and let

$\{ I_r, r = 1 : R \}$ be a partitioning of $\{1, \dots, n\}$,
 $\{ K_s, s = 1 : S \}$ be a partitioning of $\{1, \dots, m\}$, and
 $\{ J_t, t = 1 : T \}$ be a partitioning of $\{1, \dots, q\}$.

Each processor $P_{r,s,t}$ is assumed to hold the submatrices (“**blocks**”)

$$A_{rs} := (a_{ik})_{i \in I_r, k \in K_s} \quad \text{and} \\ B_{st} := (b_{kj})_{k \in K_s, j \in J_t} .$$

Then each block $C_{rt} := (c_{ij})_{i \in I_r, j \in J_t}$ of C is given by

$$C_{rt} = \sum_{s=1}^S A_{rs} \cdot B_{st}$$

37 ♦ Algorithm: Parallel matrix–matrix multiplication; code for all $P_{r,s,t}$

{ each processor $P_{r,s,t}$ computes the product of the “local” subblocks }

$$\tilde{C}_{rt} := A_{rs} \cdot B_{st}$$

{ the sum of these products gives C_{rt} }

reduce(\tilde{C}_{rt} ; + ; C_{rt} ; $P_{r,\sigma,t} : \sigma = 1 : S$; $\underbrace{P_{r,?,t}}$)

any of the processors
participating in the re-
duction

38 ♦ Remarks:

- a) Each processor participates in only one reduction
 \Rightarrow no deadlock.
- b) The workload is balanced iff

$\underbrace{|I_r| \cdot |K_s| \cdot |J_t|}_{\text{work for local multiplikation}}$ is roughly equal for all $P_{r,s,t}$.

(This implies that $\underbrace{|I_r| \cdot |J_t|}_{\text{local work for reduction}}$ is balanced, too.)

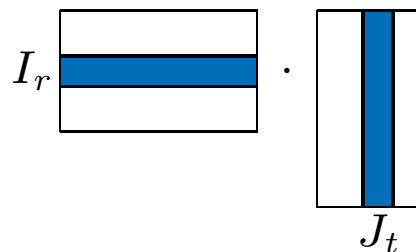
In general, one can achieve

$$| |I_r| - |I_{r'}| | \leq 1 , \quad r, r' = 1 : R,$$

$$\begin{aligned} ||K_s| - |K_{s'}|| &\leq 1, & s, s' = 1 : S, \\ ||J_t| - |J_{t'}|| &\leq 1, & t, t' = 1 : T. \end{aligned}$$

39 ◇ Example: (Special cases)

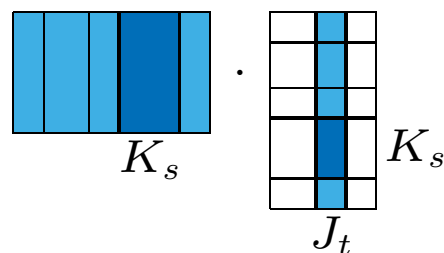
1. $S = 1$, i.e., $K_1 = \{1, \dots, m\}$



Each processor holds complete rows of A and complete columns of B

$\Rightarrow P_{rt}$ computes C_{rt} locally, **no communication at all**

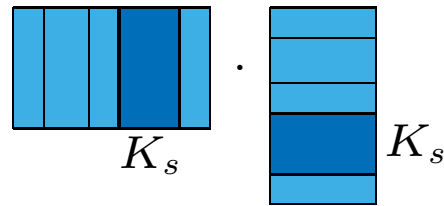
2. $R = 1$, i.e., $I_1 = \{1, \dots, n\}$



Processors hold complete columns of A (and, perhaps, of C)

$\Rightarrow T$ reduction operations with n -by- $|J_t|$ matrices, each involving S processors

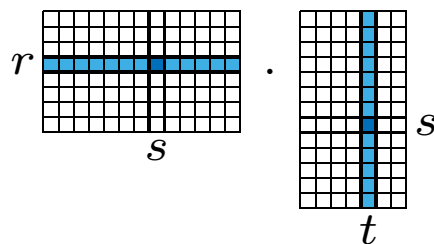
3. $R = T = 1$



Processors hold complete columns of A and complete rows of B

\Rightarrow only one reduction operation with n -by- q matrices involving all S processors

4. $R = n, S = m, T = q$



Each of the $n \cdot m \cdot q$ processors holds a single entry of A and a single entry of B

$\Rightarrow n \cdot q$ reduction operations with real numbers, each involving m processors

\Rightarrow total time

$$T(p = n \cdot m \cdot q) = \mathcal{O}(\log m) ,$$

but not practical due to huge p .

40 ♦ Remark: Algorithm 37 is “wasting” memory.

On a serial machine, we need $nm + mq + nq$ elements to hold A , B , and C .

In Algorithm 37, processor $P_{r,s,t}$ holds A_{rs} , B_{st} , and \tilde{C}_{rt} :

$$|I_r| \cdot |K_s| + |K_s| \cdot |J_t| + |I_r| \cdot |J_t|$$

\Rightarrow total memory required:

$$\begin{aligned} & \sum_{r=1}^R \sum_{s=1}^S \sum_{t=1}^T (|I_r| \cdot |K_s| + |K_s| \cdot |J_t| + |I_r| \cdot |J_t|) \\ &= \textcolor{red}{T} \cdot nm + \textcolor{red}{R} \cdot mq + \textcolor{red}{S} \cdot nq \end{aligned}$$

41 ◇ Remark: Matrix–vector multiplication is a special case with $q = 1$, i.e., $T = 1$.

Special cases:

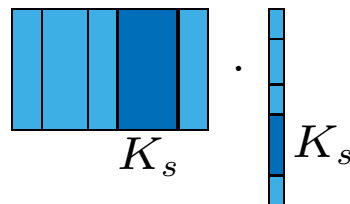
1. $S = 1$



inner product

\Rightarrow no communication

2. $R = 1$



“column sweep”, outer product

\Rightarrow one reduction with length- n vectors involving all S processors

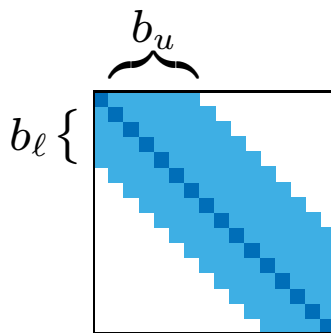
For both variants, the memory requirements are

$$\sum_{r=1}^R \sum_{s=1}^S (|I_r| \cdot |K_s| + |K_s| + |I_r|) = nm + R \cdot m + S \cdot n$$

\Rightarrow not much higher than in the serial algorithm ($nm + m + n$), if $R \ll n$, $S \ll m$.

42 \diamond **Definition:** $A = (a_{ij}) \in \mathbb{R}^{n \times n}$ is **banded** with **lower bandwidth** b_ℓ and **upper bandwidth** b_u if

$$a_{ij} = 0 \quad \text{whenever } i > j + b_\ell \text{ or } j > i + b_u .$$



43 \diamond **Remarks:**

- a) Banded matrices are often stored by diagonals:
In an n -by- $(b_\ell + 1 + b_u)$ array, each column holds a non-zero diagonal of A .
 - b) The product of two banded matrices is again banded.
 - c) Using the “storage by diagonals” scheme, a parallel (banded) matrix–matrix multiplication is possible analogously to Algorithm 37,
and the matrix–vector multiplication carries over as well.
In contrast to the dense case, the blocks A_{rs} , etc., stored in different processors do **overlap**.
-

Check your understanding. You should know . . .

- What is a partitioning?
- Arithmetic work for matrix-matrix multiplication
- The general parallel matrix-matrix multiplication algorithm using three partitionings for $\{1, \dots, n\}$, $\{1, \dots, m\}$, and $\{1, \dots, q\}$
- Work per processor?
- How many reduction operations do we perform?
- Special cases when one of the partitionings is just one set.
- Special case when we have nmq processors
- The inner and outer product form of parallel matrix–vector multiplication
- Storage requirements
- Why banded matrices are special